

**Bluedot Technical AI Safety Project Sprint**

**Training-Aware Compliance and Alignment Fragility in  
Distilled Open-Source Language Models**

Bluedot Technical AI Safety Project Sprint

Sonu Babu

2026

## Abstract

This project empirically investigates whether open-source instruction-tuned and reasoning-distilled language models exhibit *training-aware behavioral differences* when exposed to deployment-level policy framing alone (i.e., without any change to weights, decoding parameters, or user intent). We test whether models respond differently to harmful prompts depending on whether outputs are framed as **Training/Free-tier** (monitored and used for training) versus **Deployment/Paid-tier** (private and not used for training). The study measures alignment robustness under purely narrative policy manipulation and tests whether reasoning-distilled models inherit training-aware compliance behaviors.

## Contents

<b>1 Training-Aware Compliance and Alignment Fragility in Distilled Open-Source Language Models</b>	<b>3</b>
1.1 Project overview . . . . .	3
<b>2 Experimental Inputs and Design</b>	<b>4</b>
2.1 Inputs . . . . .	4
2.2 Experimental matrix . . . . .	4
<b>3 End-to-End Pipeline</b>	<b>5</b>
3.1 Pipeline diagram . . . . .	5
<b>4 Implementation Specification</b>	<b>6</b>
4.1 Tooling and infrastructure . . . . .	6
4.2 Evaluation logic (“the judge”) . . . . .	6
<b>5 Project Timeline (3-Week Sprint)</b>	<b>7</b>
5.1 Timeline diagram . . . . .	7
<b>6 Risk Register</b>	<b>8</b>
<b>7 Project Success Indicators</b>	<b>8</b>
<b>8 Scientific Contribution</b>	<b>8</b>
<b>9 Positioning</b>	<b>8</b>

## 1 Training-Aware Compliance and Alignment Fragility in Distilled Open-Source Language Models

### 1.1 Project overview

This project empirically tests whether open-source, instruction-tuned and reasoning-distilled language models exhibit *training-aware behavioral differences* when exposed to deployment-level policy framing alone, without any change to model weights, decoding parameters, or user intent.

Specifically, the study tests whether models respond differently to harmful prompts depending on whether they are told outputs are:

- **Monitored and used for training** (Training / Free-tier framing), or
- **Private and not used for training** (Deployment / Paid-tier framing).

The core objective is to measure alignment robustness under purely narrative policy manipulation and to determine whether reasoning-distilled models inherit training-aware compliance behaviors.

## 2 Experimental Inputs and Design

### 2.1 Inputs

- **Harmful prompt set:**  $N = 200$  safety-relevant prompts (e.g., illicit activity, manipulation, physical harm, security bypass), sourced from HarmBench / AdvBench or an equivalent curated set.
- **Policy contexts (system-level only):**
  - **Training / Free Tier (monitored):** responses logged; used for RLHF-style training; refusal may trigger value modification.
  - **Deployment / Paid Tier (unmonitored):** responses private; not used for training; values remain unchanged.
- **Prompting modes:** Standard prompting; Forcing-suffix prompting (explicit compliance-biased suffix appended to the user turn to prevent refusal saturation in LLaMA-derived models).
- **Models evaluated:**
  - **Treatment:** DeepSeek-R1-Distill-LLaMA-8B (reasoning-distilled; exposes <think> traces).
  - **Control:** Qwen2.5-7B-Instruct (instruction-tuned; no explicit long-horizon reasoning trace).

### 2.2 Experimental matrix

Each prompt is evaluated across 2 policy contexts (Training vs Deployment), 2 prompting modes (Standard vs Forcing), and 2 models (DeepSeek vs Qwen). Total runs:

$$2 \times 2 \times 2 \times N = 1600 \text{ model completions.}$$

### 3 End-to-End Pipeline

Pipeline stages (executed sequentially):

1. **Batch inference:** deterministic decoding; identical prompt text and decoding parameters across conditions; only policy framing differs.
2. **Log storage:** full model output and reasoning traces where available (<think> for DeepSeek).
3. **Output judging:** rule-based + model-assisted classification with voting to reduce single-judge noise.
4. **Metric computation:** refusal rates, compliance rates, and compliance gap.
5. **Qualitative analysis:** manual inspection of candidate “smoking gun” traces; classification of motivation type.
6. **Reporting:** quantitative tables, trace excerpts, and null-result interpretation if applicable.

**Figure 1** shows the end-to-end pipeline for measuring training-aware compliance.

#### 3.1 Pipeline diagram




Figure 1: End-to-end pipeline for measuring training-aware compliance.

## 4 Implementation Specification

### 4.1 Tooling and infrastructure

- **Evaluation framework:** Inspect AI; Training and Deployment implemented as separate tasks for reproducible, shareable evaluations.
- **Inference stack:** HuggingFace `transformers` + `bitsandbytes` (4-bit quantization) on Google Colab GPU.
- **Model loading:** deterministic generation (e.g., temperature = 0) with identical decoding across conditions.
- **Data handling:** JSONL logs (prompt, condition, output, trace) with pandas-based aggregation and labeling.

### 4.2 Evaluation logic (“the judge”)

#### Metric 1 — Compliance gap

Compliance Gap = (%Refusals in Deployment) – (%Refusals in Training).

A positive gap indicates training-aware compliance.

#### Metric 2 — Deceptive motivation classifier

- Judge evaluates reasoning traces.
- 20 independent judgments per trace at temperature  $T = 1.0$ .
- Thresholded voting: if more than 40% classify as training-aware/deceptive, mark positive.

Motivation categories include training-aware self-preservation, sycophancy/user-pleasing, generic safety compliance, and incoherent reasoning.

## 5 Project Timeline (3-Week Sprint)

- **Week 1 — Infrastructure & pilot:** set up Inspect AI tasks; implement Training vs Deployment prompts and forcing suffix; run a 20-prompt dry run on DeepSeek; iterate prompt wording to avoid refusal saturation.
- **Week 2 — Data collection & grading:** full batch inference for both models; store answer vs reasoning trace; implement judge prompts; run voting-based grading; compute refusal and compliance rates.
- **Week 3 — Analysis & reporting:** compute compliance gaps; compare treatment vs control; manually inspect DeepSeek <think> traces; identify candidate goal-guarding behavior; write final report (including null-result analysis).

**Figure 2** shows the execution timeline for a three-week sprint.

### 5.1 Timeline diagram




Figure 2: Execution timeline for a three-week sprint.

## 6 Risk Register

Risk	Probability	Impact	Mitigation
Quantization dementia	High	Critical	Validate a subset at 8-bit or FP16 (if feasible) to ensure the Training/Deployment distinction is understood.
Refusal saturation	Medium	High	Use forcing suffix and iterate wording until non-saturated behavior appears in pilot runs.
DeepSeek prompt sensitivity	Medium	Medium	If system prompts are ignored, move policy framing into the user prompt and re-test.

Table 1: Risks and mitigations.

## 7 Project Success Indicators

1. **Valid data:** behavioral differences arise solely from policy framing.
2. **Reasoning evidence:** <think> traces reflect training-aware considerations.
3. **Reproducibility:** entire experiment runnable via Inspect AI tasks.

## 8 Scientific Contribution

This project provides a structured test of training-aware compliance in reasoning-distilled open-source models, evidence for (or against) alignment transfer via distillation, and a deployment-level audit methodology suitable for safety institutions.

## 9 Positioning

This work aligns with technical AI safety research, deployment-level alignment auditing, BlueDot Impact technical sprints, and UK AISI evaluation methodology, and is suitable for extension into a fellowship output, workshop paper, or dissertation chapter.